

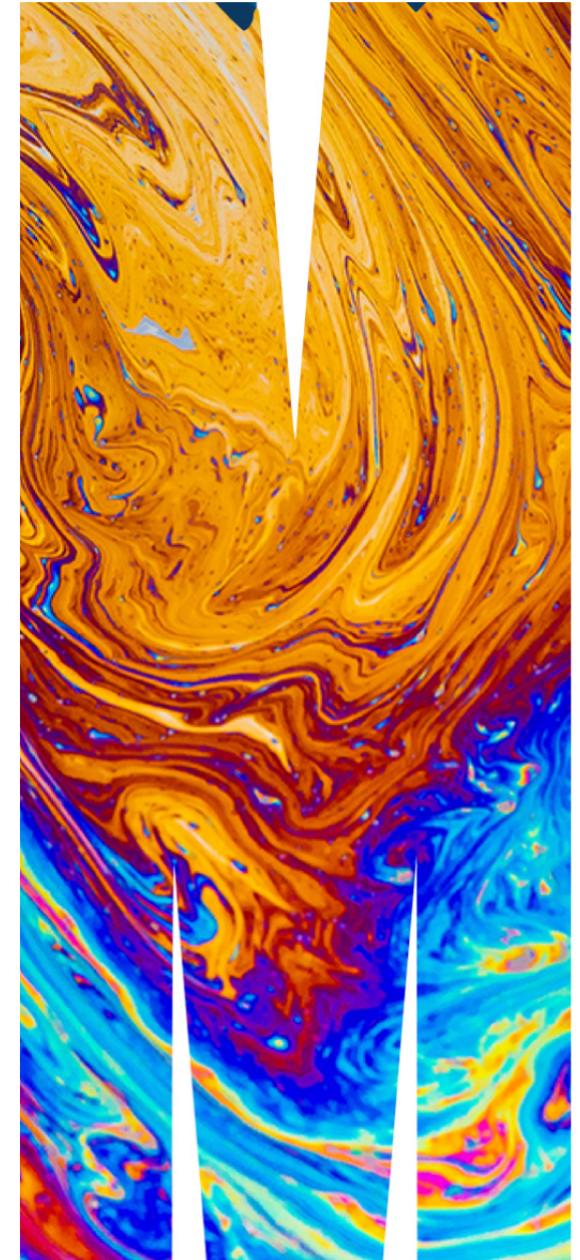
ETC5521: Exploratory Data Analysis

Sculpting data using models, checking assumptions, co-dependency and performing diagnostics

Lecturer: *Di Cook*

✉ ETC5521.Clayton-x@monash.edu

CALENDAR Week 11 - Session 1



Models help focus on the structure



before focus



after focus, we can see it's a rare Eurasian Hoopoe.

Parametric regression

Parametric regression

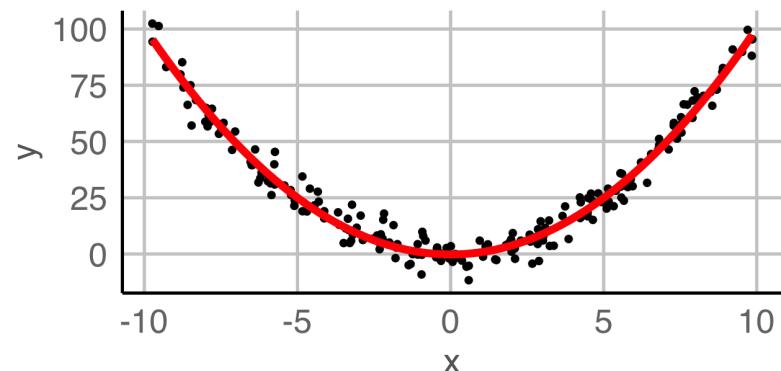
- **Parametric** means that the researcher or analyst assumes in advance that the data fits some type of distribution (e.g. the normal distribution).
- E.g. one may assume that

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \epsilon_i,$$

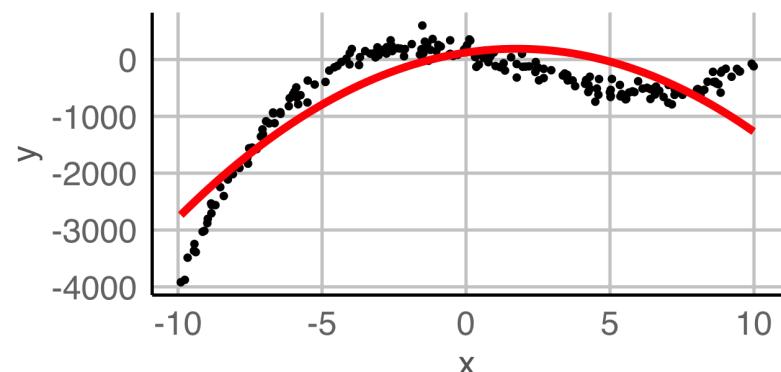
where $\epsilon_i \sim NID(0, \sigma^2)$ for $i = 1, \dots, n$,

- red = to estimate
- blue = observed
- Because some type of distribution is assumed in advance, parametric fitting can lead to fitting a smooth curve that misrepresents the data.

Examples



Assuming a quadratic fit:



Simulating data from parametric models

- Say a model is

$$y = x^2 + e, \quad e \sim N(0, 2^2).$$

- Then we have

$$y | x \sim N(x^2, 2^2).$$

Simulating data from parametric models

- Say a model is

$$y = x^2 + e, \quad e \sim N(0, 2^2).$$

- Then we have

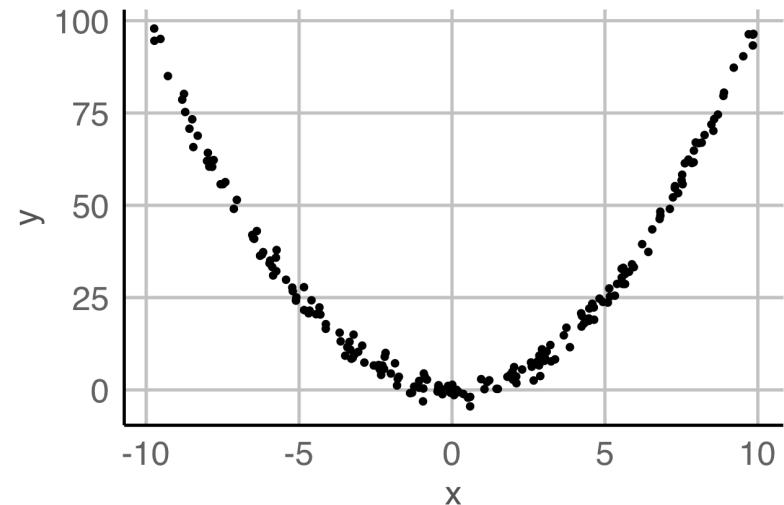
$$y | x \sim N(x^2, 2^2).$$

- Let's draw 200 observations from this model.
- Suppose that $x \in (-10, 10)$ and that we have uniform coverage over the support.
- The response y is generated as per above model.

```
set.seed(1)
df <- tibble(id = 1:200) %>%
  mutate(x = runif(n(), -10, 10),
        y = x^2 + rnorm(n(), 0, 2))
```

Plotting this:

```
ggplot(df, aes(x, y)) +
  geom_point()
```



Logistic regression

Logistic regression

- Not all parametric models assume Normally distributed errors nor continuous responses.
- Logistic regression models the relationship between a set of explanatory variables (x_{i1}, \dots, x_{ik}) and a set of **binary outcomes** Y_i for $i = 1, \dots, n$.
- We assume that $Y_i \sim B(1, p_i) \equiv \text{Bernoulli}(p_i)$ and the model is given by

$$\text{logit}(p_i) = \ln \left(\frac{p_i}{1 - p_i} \right) = \beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik}.$$

- Taking the exponential of both sides and rearranging we get

$$p_i = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik})}}.$$

- The function $f(p) = \ln \left(\frac{p}{1 - p} \right)$ is called the **logit** function, continuous with range $(-\infty, \infty)$, and if p is the probability of an event, $f(p)$ is the log of the odds.

Representation of data for binary outcomes

Data:

```
mock_df
```

```
## # A tibble: 18 × 5
##   Patient Smoker Sex    Cancer CancerBinary
##   <fct>   <fct>  <fct>  <fct>      <dbl>
## 1 1       Yes    Female  No        0
## 2 2       Yes    Male    No        0
## 3 3       No     Female  Yes       1
## 4 4       Yes    Male    No        0
## 5 5       Yes    Female  Yes       1
## 6 6       No     Female  No        0
## 7 7       Yes    Female  Yes       1
## 8 8       No     Female  No        0
## 9 9       No     Female  No        0
## 10 10    No     Male    No        0
## 11 11    Yes    Male    No        0
## 12 12    Yes    Female  Yes       1
## 13 13    Yes    Male    No        0
## 14 14    Yes    Female  No        0
## 15 15    No     Male    Yes      1
## 16 16    No     Female  Yes       1
## 17 17    No     Male    No        0
## 18 18    No     Male    Yes      1
```

Summarised data:

```
mock_sumdf
```

```
## # A tibble: 4 × 4
## # Groups:   Smoker [2]
##   Smoker Sex    Cancer Total
##   <fct>  <fct>  <int> <int>
## 1 No     Female  2     5
## 2 No     Male    2     4
## 3 Yes    Female  3     5
## 4 Yes    Male    0     4
```

- The summarised data here give the same information as the original data, except you lost the patient number
- Note the sample size, n, is larger than the number of rows in the summarised data

Logistic regression in R

- Fitting logistic regression models in R depend on the form of input data

```
glm(Cancer ~ Smoker + Sex,  
    family = binomial(link = "logit"),  
    data = mock_df)  
  
##  
## Call: glm(formula = Cancer ~ Smoker + Sex, family = bi  
##      data = mock_df)  
##  
## Coefficients:  
## (Intercept)  SmokerYes     SexMale  
##      0.2517     -0.5034     -1.1145  
##  
## Degrees of Freedom: 17 Total (i.e. Null); 15 Residual  
## Null Deviance:      24.06  
## Residual Deviance: 22.61      AIC: 28.61
```

```
glm(cbind(Cancer, Total - Cancer) ~ Smoker + Sex,  
    family = binomial(link = "logit"),  
    data = mock_sumdf)  
  
##  
## Call: glm(formula = cbind(Cancer, Total - Cancer) ~ Smok  
##      data = mock_sumdf)  
##  
## Coefficients:  
## (Intercept)  SmokerYes     SexMale  
##      0.2517     -0.5034     -1.1145  
##  
## Degrees of Freedom: 3 Total (i.e. Null); 1 Residual  
## Null Deviance:      5.052  
## Residual Deviance: 3.604      AIC: 15.82
```

Simulating from a logistic regression model Part 1

- Let's suppose that the probability of having cancer are the following:
 - 0.075 for women smokers
 - 0.045 for men smokers
 - 0.005 for women non-smokers
 - 0.003 for men non-smokers
- We'll sample 500 people for each group
- Remember that under the logistic regression model, we assumed that $Y_i \sim B(1, p_i)$

```
df <- tibble(id = 1:2000) %>%
  mutate(Smoker = rep(c("Yes", "No"), each = n() / 2),
         Sex = rep(c("Female", "Male"), times = n() / 2)) %>%
  rowwise() %>%
  mutate(CancerBinary =
    case_when(Smoker=="Yes" & Sex=="Female" ~ rbinom(1, 1, 0.075),
              Smoker=="Yes" & Sex=="Male" ~ rbinom(1, 1, 0.045),
              Smoker=="No" & Sex=="Female" ~ rbinom(1, 1, 0.005),
              Smoker=="No" & Sex=="Male" ~ rbinom(1, 1, 0.003)),
    Cancer = ifelse(CancerBinary, "Yes", "No"))

df %>%
  filter(Cancer=="Yes")

## # A tibble: 53 × 5
## # Rowwise:
##       id Smoker Sex   CancerBinary Cancer
##   <int> <chr>  <chr>      <int> <chr>
## 1     1 Yes    Female      1 Yes
## 2     2 Yes    Male       1 Yes
## 3     3 Yes    Female      1 Yes
## 4     4 Yes    Female      1 Yes
## 5     5 Yes    Female      1 Yes
## 6     6 Yes    Male       1 Yes
## 7     7 Yes    Female      1 Yes
## 8     8 Yes    Male       1 Yes
## ...  9 245 Yes   Female      1 Yes
```

Simulating from a logistic regression model Part 2

- At times, you may want to **simulate the summary data directly** instead of the individual data
- Recall that if $Y_i \sim B(1, p)$ for $i = 1, \dots, k$ and Y_i 's are independent,

$$S = Y_1 + Y_2 + \dots + Y_k \sim B(k, p)$$

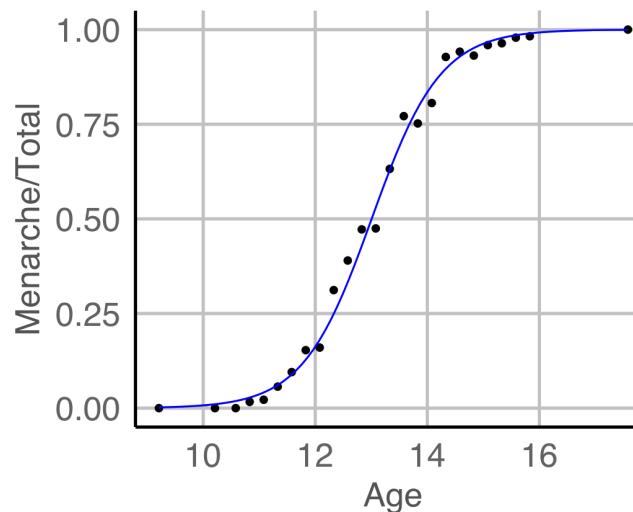
```
expand_grid(Smoker = c("Yes", "No"), Sex = c("Female", "Male")) %>%
  rowwise() %>%
  mutate(Cancer =
    case_when(Smoker=="Yes" & Sex=="Female" ~ rbinom(1, 500, 0.07)
              Smoker=="Yes" & Sex=="Male" ~ rbinom(1, 500, 0.045)
              Smoker=="No" & Sex=="Female" ~ rbinom(1, 500, 0.005)
              Smoker=="No" & Sex=="Male" ~ rbinom(1, 500, 0.003))
    Total = 500)

## # A tibble: 4 × 4
## # Rowwise:
##   Smoker Sex     Cancer Total
##   <chr>  <chr>   <int> <dbl>
## 1 Yes    Female    35    500
## 2 Yes    Male      23    500
## 3 No     Female     0    500
## 4 No     Male      3    500
```

Case study 1 Menarche

- In 1965, the average age of 25 homogeneous groups of girls was recorded along with the number of girls who have reached menarche out of the total in each group.

 data R



Simulating data from a fitted logistic regression model Part 1

- Suppose we want to simulate from the fitted model

- We first fit the fitted model

```
fit1 <-  
  glm(cbind(Menarche, Total - Menarche) ~ Age,  
       family = "binomial",  
       data = menarche)  
(beta <- coef(fit1))  
  
## (Intercept)      Age  
## -21.226395   1.631968
```

- The fitted regression model is given as:

$$\text{logit}(\hat{p}_i) = \hat{\beta}_0 + \hat{\beta}_1 x_{i1}.$$

- Rearranging we get

$$\hat{p}_i = \frac{1}{1 + e^{-(\hat{\beta}_0 + \hat{\beta}_1 x_{i1})}}.$$

- Simulating from first principles:

```
menarche %>%  
  rowwise() %>%  
  mutate(  
    phat = 1/(1 + exp(-(beta[1] + beta[2] * Age))),  
    simMenarche = rbinom(1, Total, phat))  
  
## # A tibble: 25 × 5  
## # Rowwise:  
##       Age Total Menarche     phat simMenarche  
##     <dbl> <dbl>    <dbl>    <dbl>        <int>  
## 1 9.21    376     0 0.00203         1  
## 2 10.2     200     0 0.0103         3  
## 3 10.6     93      0 0.0187         2  
## 4 10.8     120     2 0.0279         3  
## 5 11.1     90      2 0.0413         1  
## 6 11.3     88      5 0.0609         6  
## 7 11.6     105     10 0.0888        9  
## 8 11.8     111     17 0.128        12  
## 9 12.1     100     16 0.181        17  
## 10 12.3     93     29 0.249        23  
## # i 15 more rows
```

Simulating data from a fitted logistic regression model Part 2

- An easier way to do this is to use the `simulate` function which works for many model objects in R
- Below it's simulating 3 sets of responses (i.e. counts of "success" and "failure" events) from `fit1` logistic model object

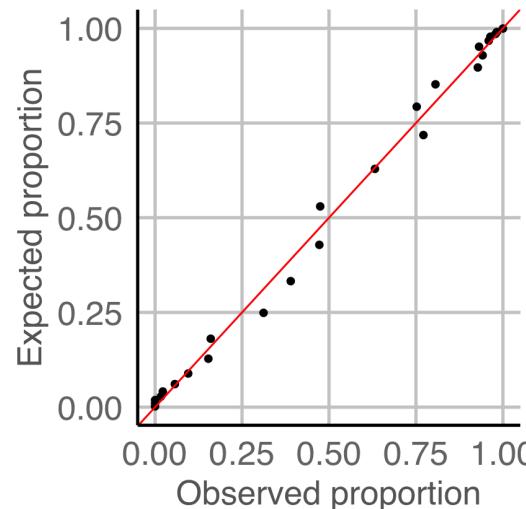
```
simulate(fit1, nsim = 3)

##      sim_1.Menarche sim_1.V2 sim_2.Menarche sim_2.V2 sim_3.Menarche sim_3.V2
## 1            0       376            0       376            0       376
## 2            2       198            1       199            0       200
## 3            4        89            0        93            3        90
## 4            4       116            2       118            6       114
## 5            8        82            5        85            3        87
## 6            6        82            3        85            6        82
## 7           14       91            7       98            5       100
## 8           13       98            14      97            16       95
## 9           21       79            18      82            20       80
## 10          25       68            21      72            27       66
## 11          47       53            32      68            33       67
## 12          37       71            43      65            41       67
## 13          47       52            54      45            59       40
```

Diagnostics for logistic regression models

- One diagnostic is to compare the observed and expected proportions under the logistic regression fit.

```
df1 <- menarche %>%
  mutate(
    pexp = 1/(1 + exp(-(beta[1] + beta[2] * Age))),
    pobs = Menarche / Total)
```



Diagnostics for logistic regression models

- Goodness-of-fit type test is used commonly to assess the fit as well.
- E.g. Hosmer–Lemeshow test, where test statistic is given as

$$H = \sum_{i=1}^r \left(\frac{(O_{1i} - E_{1g})^2}{E_{1i}} + \frac{(O_{0i} - E_{0g})^2}{E_{0i}} \right)$$

where O_{1i} (E_{1i}) and O_{0i} (E_{0i}) are observed (expected) frequencies for successful and non-successful events for group i , respectively.

```
vcdExtra::HLtest(fit1)

## Hosmer and Lemeshow Goodness-of-Fit Test
##
## Call:
## glm(formula = cbind(Menarche, Total - Menarche) ~ Age, family = "binomial",
##      data = menarche)
## ChiSquare df  P_value
## 0.1088745  8 0.9999996
```

Diagnostics for linear models

Assumptions for linear models

For $i \in \{1, \dots, n\}$,

$$Y_i = \beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik} + \epsilon_i,$$

where $\epsilon_i \sim NID(0, \sigma^2)$ or in matrix format,

$$Y = X\beta + \epsilon, \quad \epsilon \sim N(\mathbf{0}, \sigma^2 I_n)$$

where

- $\mathbf{Y} = (Y_1, \dots, Y_n)^\top$,
- $\boldsymbol{\beta} = (\beta_0, \dots, \beta_k)^\top$,
- $\boldsymbol{\epsilon} = (\epsilon_1, \dots, \epsilon_n)^\top$, and
- $\mathbf{X} = [\mathbf{1}_n \quad \mathbf{x}_1 \quad \dots \quad \mathbf{x}_k]^\top$, where
- $\mathbf{x}_j = (x_{1j}, \dots, x_{nj})^\top$ for $j \in \{1, \dots, k\}$

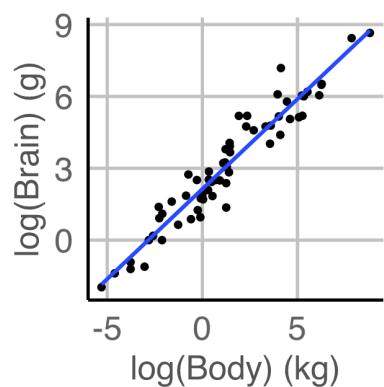
This means that we assume

1. $E(\epsilon_i) = 0$ for $i \in \{1, \dots, n\}$.
2. $\epsilon_1, \dots, \epsilon_n$ are independent.
3. $Var(\epsilon_i) = \sigma^2$ for $i \in \{1, \dots, n\}$ (i.e. homogeneity).
4. $\epsilon_1, \dots, \epsilon_n$ are normally distributed.

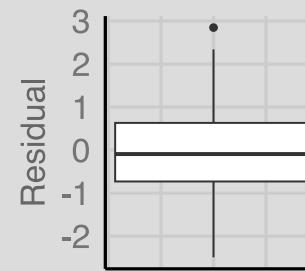
So how do we check it?

Model diagnostics for linear models

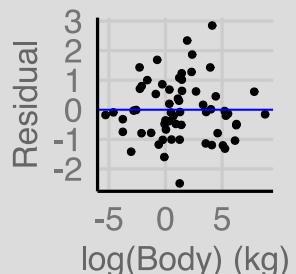
Plot Y_i vs x_i to see if there is \approx a linear relationship between Y and x .



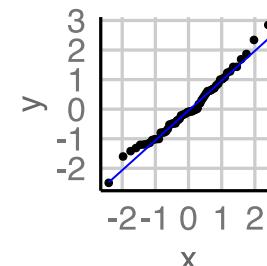
A boxplot of the residuals R_i to check for symmetry.



To check the homoscedasticity assumption, plot R_i vs x_i . There should be no obvious patterns.



A normal Q-Q plot, i.e. a plot of the ordered residuals vs $\Phi^{-1}(\frac{i}{n+1})$.



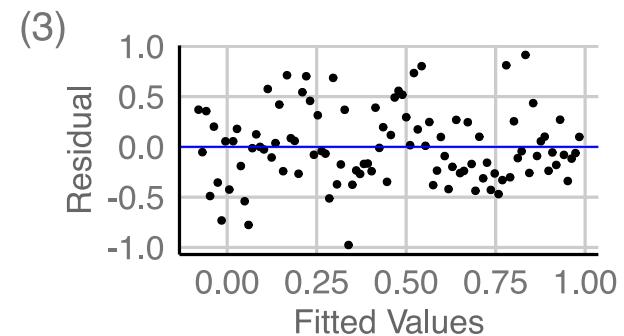
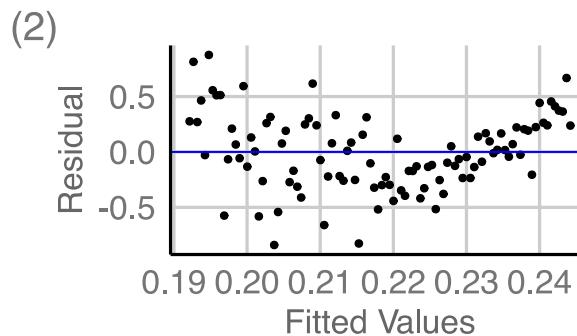
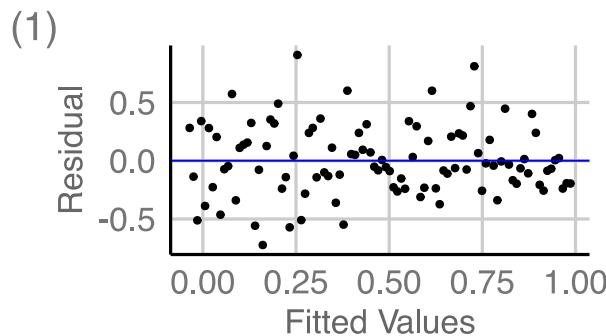
Assessing (A1) $E(\epsilon_i) = 0$ for $i = 1, \dots, n$

- It is a property of the least squares method that

$$\sum_{i=1}^n R_i = 0, \quad \text{so} \quad \bar{R}_i = 0$$

for $R_i = Y_i - \hat{Y}_i$, hence (A1) will always appear valid "overall".

- Trend in residual versus fitted values or covariate can indicate "local" failure of (A1).
- What do you conclude from the following plots?



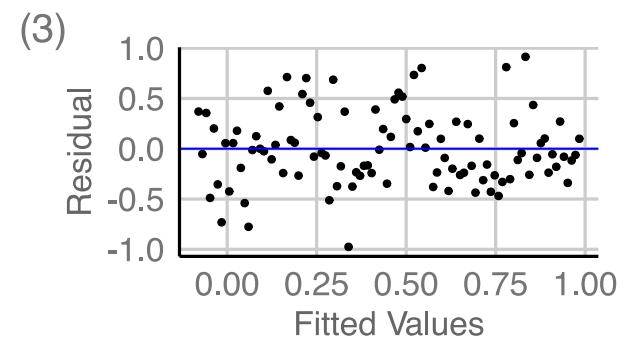
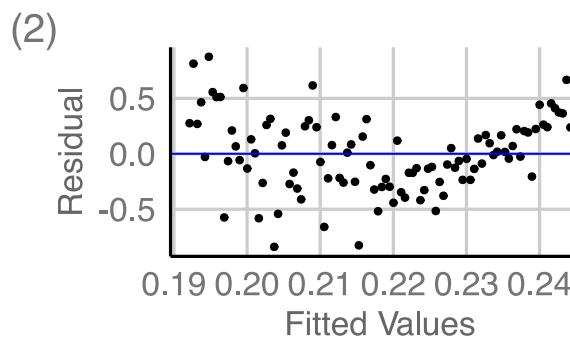
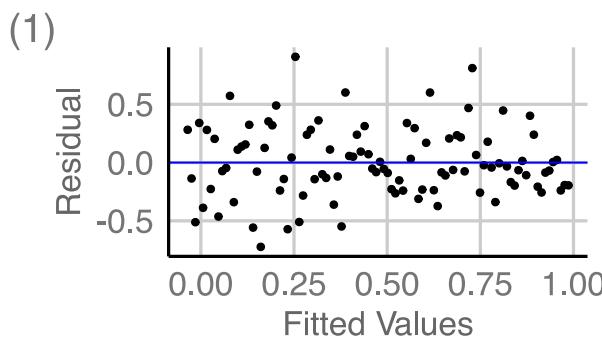
Assessing (A2)-(A3)

(A2) $\epsilon_1, \dots, \epsilon_n$ are independent

- If (A2) is correct, then residuals should appear randomly scattered about zero if plotted against fitted values or covariate.
- Long sequences of positive residuals followed by sequences of negative residuals in R_i vs x_i plot suggests that the error terms are not independent.

(A3) $\text{Var}(\epsilon_i) = \sigma^2$ for $i = 1, \dots, n$

- If (A3) holds then the spread of the residuals should be roughly the same across the fitted values or covariate.



Assessing (A4) $\epsilon_1, \dots, \epsilon_n$ are normally distributed

Q-Q Plots

- The function `qqnorm(x)` produces a Q-Q plot of the ordered vector `x` against the quantiles of the normal distribution.
- The n chosen normal quantiles $\Phi^{-1}(\frac{i}{n+1})$ are easy to calculate but more sophisticated ways exist:
 - $\frac{i}{n+1} \mapsto \frac{i-3/8}{n+1/4}$, default in `qqnorm`.
 - $\frac{i}{n+1} \mapsto \frac{i-1/3}{n+1/3}$, recommended by Hyndman and Fan (1996).

In R

```
fit <- lm(y ~ x)
```

By "hand"

```
plot(qnorm((1:n) / (n + 1)), sort(resid(fit)))
```

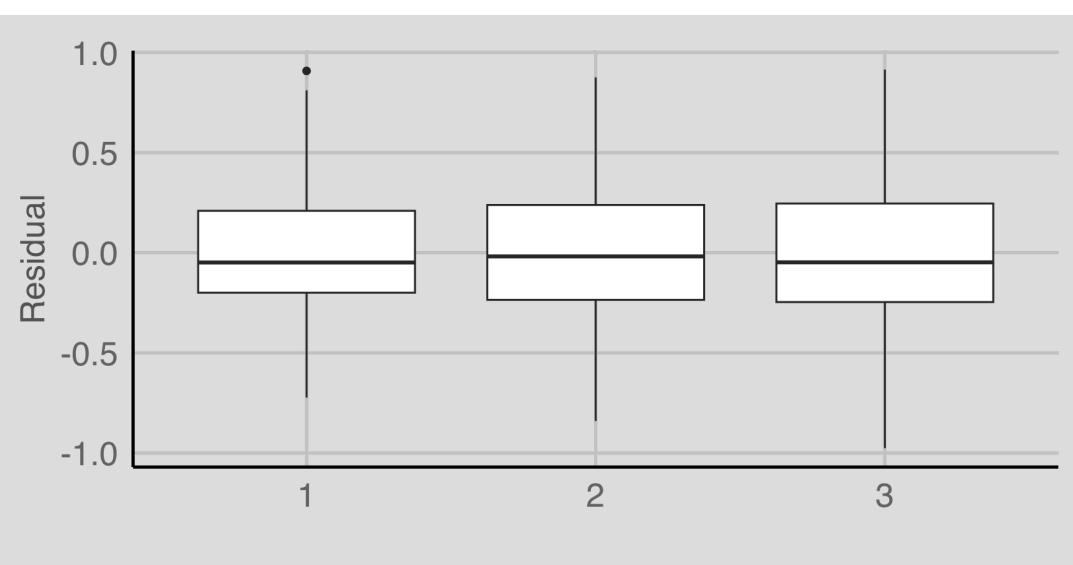
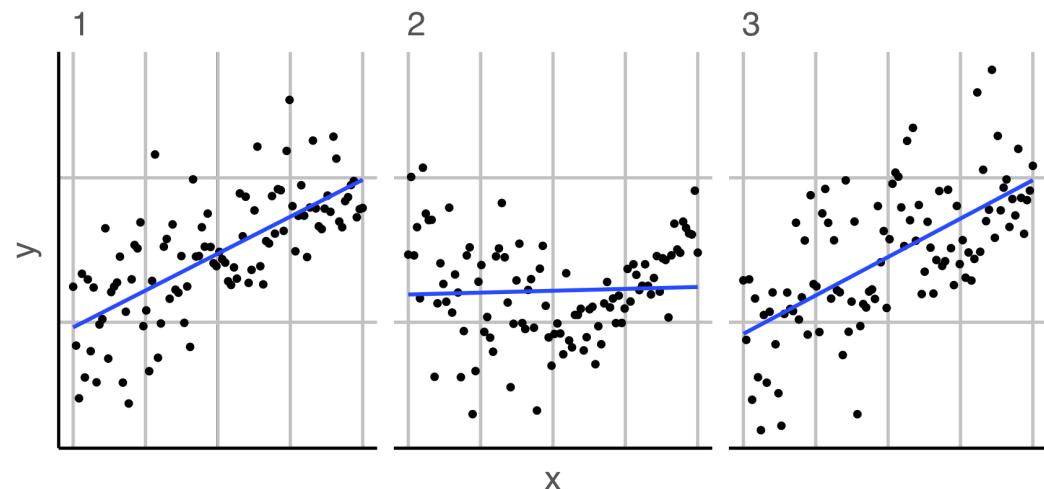
By `base`

```
qqnorm(resid(fit))  
qqline(resid(fit))
```

By `ggplot2`

```
data.frame(residual = resid(fit)) %>%  
  ggplot(aes(sample = residual)) +  
  stat_qq() + stat_qq_line(color="blue")
```

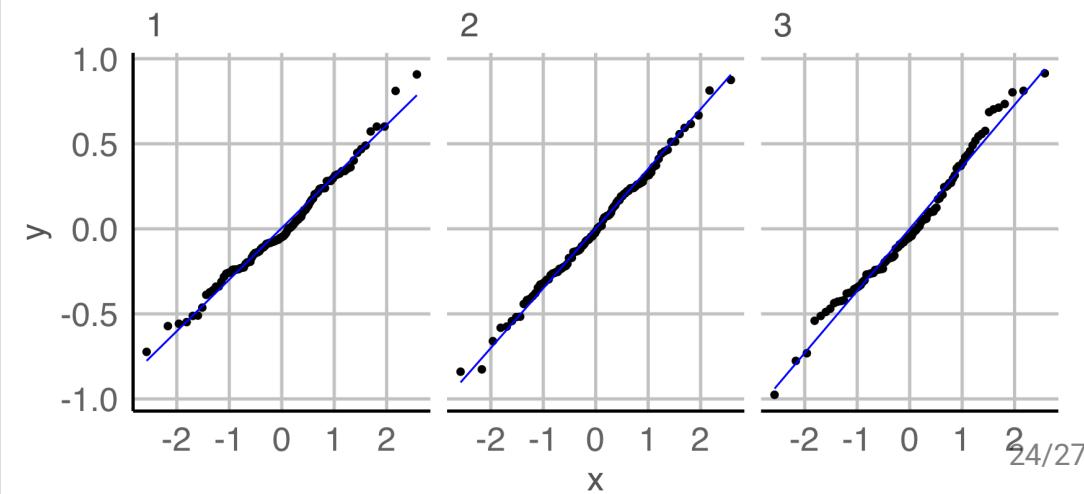
Examining simulated data



Simulation scheme

```
n <- 100
x <- seq(0, 1, length.out = n)
y1 <- x + rnorm(n) / 3 # Linear
y2 <- 3 * (x - 0.5) ^ 2 +
  c(rnorm(n / 2)/3, rnorm(n / 2)/6) # Quadratic
y3 <- -0.25 * sin(20 * x - 0.2) +
  x + rnorm(n) / 3 # Non-linear

M1 <- lm(y1 ~ x); M2 <- lm(y2 ~ x); M3 <- lm(y3 ~ x)
```



Take away messages

- Parametric models assume some distribution in advance
- Logistic models can be used to model explanatory variables with binary outcomes
- You should be able to simulate from parametric models
- You can perform basic model diagnostics
- You can use simulation to analyse model properties

Resources and Acknowledgement

- These slides were originally created by Dr Emi Tanaka, and modified by Dr Michael Lydeamore.
- Some of these slides were inspired by STAT3012 Applied Linear Models at The University of Sydney by Prof Samuel Muller
- Cook & Weisberg (1994) "An Introduction to Regression Graphics"
- Data coding using [tidyverse suite of R packages](#)
- Slides constructed with [xaringan](#), [remark.js](#), [knitr](#), and [R Markdown](#).



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](#).

Lecturer: *Di Cook*

✉ ETC5521.Clayton-x@monash.edu

CALENDAR Week 11 - Session 1

